

The Elements of Public-Key Cryptography

Keys and the Key Distribution Problem

The need to communicate in secret is as old as communication itself.¹ The time-honored approach is to scramble a message before sending it, which the receiver unscrambles on receipt. Without knowledge of the method by which the message is scrambled, unauthorized parties who intercept the message cannot unscramble it. The mechanism for scrambling, and unscrambling, is generally referred to as the *key*.²

More formally, for two (or more) parties to communicate securely over an *insecure* channel, each must possess a key that can be used to encipher and decipher messages transmitted over the insecure channel.³ This key must be possessed only by the parties authorized to send and receive the messages; otherwise an eavesdropper with possession of the key will be able to read the messages.

The requirement that only the parties authorized to participate in a secure conversation share a secret key poses a problem: How is the key distributed to the authorized parties securely; that is, without it being stolen by an unauthorized party? For this you need a *secure* channel. Transmitting the key over an insecure channel is not an option, because if the channel is insecure the key can be stolen. Neither is enciphering the key before transmitting it, since you cannot decipher it without first having the key.⁴

The most obvious and effective solution is to hand-deliver the key in advance to the party you wish to communicate with. But this is also the least efficient solution, and not at all practical in the internet age.⁵

A Clever Solution

In 1976, two Stanford University cryptographers proposed an elegant solution to the key distribution problem in a groundbreaking paper titled *New Directions in Cryptography*. This solution became, and remains to this day, the de facto standard for exchanging keys securely over insecure channels. It is commonly known as the *Diffie-Hellman key exchange protocol*.⁶

¹ King to general, via trusted courier: *Attack at dawn!*

² The key can be a mechanical device, a number, a puzzle; anything possessed by both sender and receiver that enables the sender to encipher, and the receiver to decipher, a message.

³ The quintessential example of such an insecure channel is the public internet.

⁴ The classic *chicken-and-egg* problem.

⁵ Accepting that each n -party communication requires a separate key, the number of key exchanges required for a group of n participants to communicate securely is found by following formula: $n(n-1)/2$, where n is the number of participants. For a group of 10, the number of key exchanges is $10(10-1)/2$, or 45; for 100 the number is 4,950, and so on. As the number of participants increases, the number of key exchanges increases quadratically.

⁶ Although Whitfield Diffie and Martin Hellman co-authored the paper, and their names are attributed to the protocol, Ralph Merkle's name deserves mention because it is on Merkle's ideas that Diffie-Hellman is based (see *Merkle's Puzzles*).

The Diffie-Hellman (DH) key exchange protocol is a foundational element of *public-key* cryptography;⁷ namely, that of secure key exchange over insecure channels. DH enables two (or more) parties to exchange public information over an insecure channel, and then combine it with private information to compute an identical, shared key with which to encrypt and decrypt messages sent to each other over the insecure channel. Because private information is used on either side of the channel to generate the shared key, the key cannot be observed by an eavesdropper listening on the channel.

DH can be implemented by means of any number of algorithms. Most examples in the literature cite the original implementation, which uses a *multiplicative group of integers modulo a prime number*, to demonstrate DH. This is unfortunate, because the mathematics of multiplicative groups modulo a prime are complex, and thus hinder a conceptual understanding of DH.

Simplified Diffie-Hellman

The graphic in *Figure 1* demonstrates DH using rudimentary multiplication. This example depicts a simplified implementation of DH. This should prove useful in comprehending a more complex, real-world implementation later in this paper.

Assume Alice wants to perform a secure key exchange with Bob over an insecure channel. Meanwhile, Eve observes all traffic passing between Alice and Bob, presumably for malicious purposes.⁸

Step	Alice	Eve	Bob	Calculations
1	2	->		
2		2	2	
3	3			
4	6	->		(3 x 2 = 6)
5		6	6	
6			4	
7		<-	8	(4 x 2 = 8)
8	8	8		
9	24		24	(8 x 3 = 24), (6 x 4 = 24)

Figure 1. Diffie-Hellman key exchange using multiplication.

In steps 1 and 2, Alice selects a random integer (2) and transmits it to Bob.⁹ Let's call this number the *generator*, because it will be used by Alice and Bob to generate another number; namely by multiplying the generator by a private number they will each select. Because the channel is insecure, Eve observes the value of the generator (2).

In steps 3, 4 and 5, Alice selects a random integer (3), multiplies it by the generator (2), and transmits the product of the multiplication (6) to Bob. Let's call the random number that Alice selects her *private* key, and the product of its multiplication by the generator her *public* key (public because it can be

⁷ Also known as *asymmetric*-key cryptography, public-key cryptography is based on the principle that different, though mathematically related, keys can be used for encryption and decryption; whereas traditional encryption relies on identical, or *symmetric*, keys.

⁸ The examples in this paper feature the cast of fictional characters ubiquitous in the literature: Alice, Bob and Eve.

⁹ Since computers operate with numbers, we use integers in this and all subsequent examples to represent messages and keys.

observed by Eve on the insecure channel). As expected, Eve observes Alice's public key (6). But Eve does not observe Alice's private key (3), because Alice never transmits her private key to Bob.

In steps 6, 7 and 8, Bob selects a random integer (4), multiplies it by the generator (2), and transmits the product (8) to Alice. These are Bob's private and public keys, respectively. Eve observes Bob's public key (8). Eve now knows the generator (2), Alice's public key (6) and Bob's public key (8), but neither Alice nor Bob's private keys (3 and 4).

The magic of DH appears in step 9. Alice multiplies Bob's public key (8) by her private key (3). The product of this multiplication is 24. Similarly, Bob multiplies Alice's public key (6) by his private key (4). The product of this multiplication is also 24. By using a combination of public and private information in this clever way, Alice and Bob have agreed that the number 24 will be the shared key with which to encrypt and decrypt messages between them.

Where does this leave Eve? Having only seen the value of the generator (2), Alice's public key (6) and Bob's public key (8), but neither Alice's nor Bob's private keys (3 and 4), Eve does not know by what factors Alice's and Bob's public keys were multiplied to compute the shared encryption key (24).¹⁰

Of course, in this highly simplified implementation, Eve can easily guess Alice or Bob's private keys, and with either private key she can compute the shared key and decrypt the messages.

In addition to guessing a private key, Eve must also know the *algorithm* used by Alice and Bob to compute the shared key (i.e. multiplication). We must assume that Eve knows this algorithm, and this is a perfectly reasonable assumption given that the efficacy of modern cryptography relies wholly on the secrecy of *keys*, not *algorithms*.¹¹ With knowledge of the algorithm, Eve simply divides Alice's public key (6) by the generator (2), both of which she observed on the insecure channel, to derive Alice's private key (3).

By using division, Eve performs the *inverse* of the multiplication Alice used to generate her public key.¹² Similarly, Eve can divide Bob's public key (8) by the generator (2) to derive Bob's private key (4). The important point is that with either Alice or Bob's private key, Eve can compute the shared key and use it to decrypt all messages exchanged between Alice and Bob.

It should not be surprising that a DH implementation using multiplication to generate shared keys is badly flawed. But the object of this example is not to demonstrate an effective DH implementation, but to demonstrate how Alice and Bob can generate an identical, secret key using a combination of public and private information.

¹⁰ It is useful to point out here that Alice and Bob could have selected any private key (besides 3 and 4, respectively) and the effect in step 9 would have been the same: i.e. identical secret keys.

¹¹ This fact is formalized in Kerckhoffs's principle, proposed by Auguste Kerckhoffs in 1883, which turned several millennia of cryptographic orthodoxy on its head. Kerckhoffs stated that, "A cryptosystem should be secure even if everything about the system, except the key, is public knowledge". Prior to this, the efficacy of a cipher was believed to be based on the secrecy of its algorithm. One important implication of Kerckhoffs's principle is that a cipher whose algorithm is widely-known will invite attacks, and that this is desirable because very smart people know they will become famous if they find a way to defeat it. It should not be surprising that the best cryptosystems in the world are those that have defied successful attacks over a long period of time.

¹² Division is the inverse of multiplication, just as subtraction is the inverse of addition.

For an effective DH implementation, we need to make the task of guessing Alice and Bob's private keys more difficult for Eve.

DH With Exponentiation

Let's look at a second example using slightly more sophisticated math to make Eve's task more difficult.

Step	Alice	Eve	Bob	Calculations
1	2	->		
2		2	2	
3	3			
4	8	->		$(2^3 = 8)$
5		8	8	
6			4	
7		<-	16	$(2^4 = 16)$
8	16	16		
9	4096		4096	$(16^3 = 4096), (8^4 = 4096)$

Figure 2. Diffie-Hellman key exchange using exponentiation.

Instead of using multiplication to generate public keys, this time Alice and Bob use *exponentiation* (exponentiation in the figure is denoted by the '^' symbol; as in $10^3 = 10 \times 10 \times 10 = 1000$). Except for the calculations, all the steps are the same as in the previous example, so it's not necessary to repeat them here. What is different this time is (a) the algorithm used to compute the keys—exponentiation versus multiplication—and (b) the public parameters observed by Eve.

As before, Eve observes the generator (2), Alice's public key (8) and Bob's public key (16). With this information Eve must be able to compute the shared secret key (4096) to break the encryption. Because Eve knows the algorithm, she knows that Alice raised the generator (2) to the power of some exponent to compute her public key (8). To find that exponent, Eve must solve for y in the equation $x^y = z$, where x and z are known.¹³

Solving for y in this equation is known as taking the logarithm of z . Taking a logarithm is the inverse of exponentiation, just as division is the inverse of multiplication, which Eve used in the previous example to break the encryption. For small values of z (8 and 16 in the present example) solving for y is trivial; it simply requires trying consecutive exponents until the right answer is found. The requisite number of tries to find Alice's private key in the present example is three. For larger values of z , however, the complexity of Eve's task grows proportionately with the size of z . Contrast this with the implementation using multiplication, in which the complexity of the task is constant relative to the size of the public key.

As the value of z increases, solving for its logarithm requires more tries, and if z is large enough it becomes computationally infeasible. With this we get closer to an effective DH implementation, but we're not quite there yet.

DH With Modular Exponentiation

Modular exponentiation brings us finally to the realm of real-world DH.

¹³ In the present example, what is the value of y in the equation $2^y = 8$?

Step	Alice	Eve	Bob	Calculations
1	3			
2	7	->		
3		(3,7)	3	
4			7	
5	3			
6	6	->		$(3^3 = 6 \text{ mod } 7)$
7		6	6	
8			4	
9		<-	4	$(3^4 = 4 \text{ mod } 7)$
10	4	4		
11	1		1	$(4^3 = 1 \text{ mod } 7), (6^4 = 1 \text{ mod } 7)$

Figure 3. Diffie-Hellman key exchange using modular exponentiation.

Modular exponentiation is the same as exponentiation, but with an additional step. This additional step is called *taking a modulus*. This is done with the modulo operation.¹⁴ If you compare the graphic in this example with the one that uses exponentiation only (Figure 2), you will find the only difference is that in this one the modulo step is added to all the calculations.

The modulo operation requires an *operand*—namely the number by which to divide in order to find a remainder. We call this operand the *divisor*. This explains why Alice transmits two numbers to Bob (steps 1 and 2) instead of one, like she did in the previous examples.

As before, Alice transmits the generator (3). But she also transmits the divisor (7) that will be used to compute remainders. Note that the divisor too can be observed by Eve; it is a public parameter.¹⁵

The modular exponentiation of real-world DH leads to a very interesting property of the generated keys. Compare the values of the public and shared keys in this example (6, 4 and 1) to those of the example in Figure 2 (8, 16 and 4096). With modular exponentiation the keys are smaller; notably, they are confined to the set of positive integers starting at one, and ending at one less than the divisor.¹⁶

The graphic in Figure 4 should clarify why this is so. The values in the *Power* column are the result of raising the generator (3) to the power of the values in the *Exponent* column, and the values in the *Mod* column are the result of performing the modulo operation on the corresponding value in the *Power* column. Note that all our keys (6, 4 and 1) are present in the *Mod* column.

¹⁴ A modulo operation simply finds the remainder after division of two numbers. For example, $7 \text{ mod } 3 = 1$, because 7 divided by 3 equals 2, leaving 1 left over.

¹⁵ Public parameters in public-key cryptography are often also referred to in the literature as *domain* parameters.

¹⁶ Exponentiation of a generator g modulo p , where g is greater than 1 and p is a prime number, guarantees result will be within the set 1 to $p - 1$; in the present example, where $g = 3$ and $p = 7$, this set contains the integers 1, 2, 3, 4, 5, and 6.

Exponent	Power	Mod	Calculations
1	3	3	$(3^1 \bmod 7 = 3)$
2	9	2	$(3^2 \bmod 7 = 2)$
3	27	6	$(3^3 \bmod 7 = 6)$
4	81	4	...
5	243	5	
6	729	1	
7	2187	3	
8	6561	2	
9	19683	6	
10	59049	4	
11	177147	5	
12	531441	1	

Figure 4. The multiplicative group of integers modulo 7.

The keys exhibit two more interesting properties: First, if you read straight down the *Mod* column, you find that the sequence of remainders repeats after a while (3, 2, 6, 4, 5, 1, 3, 2, 6, 4, 5, 1). Second, each repeating sequence contains every integer in the set 1 to $p - 1$.¹⁷

To keep things simple, it will suffice to keep the following rule in mind: Given a carefully chosen generator, and a prime divisor, we can generate keys with the aforementioned properties. And it is these properties that are required for an effective DH implementation.¹⁸

Recall from the previous example that Eve had to solve for the logarithm of z in the equation $x^y = z$ to find the shared key. In the finite group of integers modulo p , this is qualified a bit; it is called the *discrete logarithm problem*, which is often abbreviated in the literature to DLP.

There is no *known*, efficient algorithm for solving the DLP. Consecutive integers must be tried in the equation $x^y = z \bmod p$, where the values of x , z and p are known, until the correct y is found. If z is large enough, the number of tries becomes too computationally expensive to be feasible for an attacker. The efficacy of DH is based, at least in part, on the difficulty of solving the DLP.¹⁹

Toward a More Complete Cryptosystem

Whereas DH solves the problem of secure key exchange, what of encryption itself; the principal use case for cryptography? Besides secure key exchange, encryption is one of the key elements of the public-key

¹⁷ These properties are described formally in the language of abstract algebra; specifically, number theory and multiplicative groups modulo p , where p is a prime number.

¹⁸ A carefully chosen generator is one which generates the entire group of integers in the range 1 to $p - 1$, where p is the prime modulus. Any generator that fulfills this property is called a *primitive root*, and the group it produces a *cyclic* group. The rules of multiplicative groups modulo p guarantee that at least one integer in the group 1 to $p - 1$ is a primitive root. In the current example, 3 is a primitive root of the cyclic group of integers modulo 7. In real-world DH, the modulus should be a very large, randomly-chosen prime number. We use an artificially small value (7) for the divisor in the example to demonstrate the concepts.

¹⁹ It is possible there is some other, as yet unknown (or at least unpublished), way to break DH; that is, besides solving the DLP efficiently. Until or unless such a method is found, no distinction is made between the DLP and the so-called *Diffie-Hellman problem*.

cryptosystem described, but not solved, by Diffie and Hellman in their 1976 paper.²⁰ For the answer to that question the crypto community would have to wait another two years.

Public-Key Encryption

The public-key cryptosystem conceived by Diffie and Hellman in *New Directions* consists of three distinct but interrelated elements: *secure key exchange*, *encryption* and *digital signatures*. But the paper only presented an actual implementation for secure key exchange.

While doing research at MIT in 1978, the cryptographers Ronald Rivest, Adi Shamir and Leonard Adelman published a paper called *A Method for Obtaining Digital Signatures and Public-Key Cryptosystems*. In it they picked up where Diffie and Hellman left off, presenting practical implementations for public-key encryption and digital signatures.

To this day, the contributions of Rivest, Shamir and Adelman form the basis the most widely used and battle tested public-key cryptosystem in the world, known simply by the initials of the surnames of its authors. RSA, or some variation of it, secures literally trillions of dollars of electronic funds transfers *daily*.

To understand RSA encryption, we'll use the same approach we took with DH; by starting with a primitive example before advancing to a more realistic one.

Substitution Cipher

Figure 5 depicts a message exchange between Alice and Bob using a simple substitution cipher. This cipher employs the now-familiar mathematics of modular exponentiation.

Here, Bob wants to transmit a private message to Alice over an insecure channel. In order to prevent Eve from reading the message, Bob encrypts the message prior to sending it to Alice. On receipt, Alice decrypts the message by inverting it to its original, unencrypted form.

Step	Alice	Eve	Bob	Calculations
1	11			
2	26	->		
3		(11,26)	11	
4			26	
5			3	
6		<-	7	$(3 \times 11 = 7 \text{ mod } 26)$
7	7	7		
8	19			
9	3			$(7 \times 19 = 3 \text{ mod } 26)$

Figure 5. Encryption and decryption using a substitution cipher.

²⁰ There was a third element, *digital signatures*, which would enable the sender of a message to prove both that it originated from the sender, and that its contents were unaltered.

In steps 1, 2, 3 and 4, Alice selects a public encryption key (11) and a divisor (26), and transmits both to Bob.²¹ Eve observes the values of both the public encryption key and the divisor. We call the encryption key *public* because it can be observed by an eavesdropper.

In steps 5, 6 and 7, Bob creates a plaintext message (3),²² multiplies it by Alice's public encryption key (11), and takes the modulus of the product to compute the ciphertext ($3 \times 11 = 7 \bmod 26$).²³ Bob transmits the ciphertext (7) to Alice, which Eve also observes.

In steps 8 and 9, Alice decrypts the ciphertext. She multiplies the ciphertext (7) she received from Bob by her private decryption key (19), and takes the modulus of the product to arrive back at the plaintext ($7 \times 19 = 3 \bmod 26$). We call Alice's decryption key *private* because it is never transmitted to Bob, and therefore cannot be observed by Eve.

This is brilliant, but where does Alice's private decryption key (19) come from? The graphic in *Figure 6* gives us the answer.

Keys (e, d)	Divisor n	Inverse?		
		$\gcd(e, n)$	$d \times e = 1 \bmod n$	Calculations
1	26	1	11	$(11 \times 1 = 11 \bmod 26)$
2	26	2	22	$(11 \times 2 = 22 \bmod 26)$
3	26	1	7	$(11 \times 3 = 7 \bmod 26)$
4	26	2	18	$(11 \times 4 = 18 \bmod 26)$
5	26	1	3	...
6	26	2	14	
7	26	1	25	
8	26	2	10	
9	26	1	21	
10	26	2	6	
11	26	1	17	
12	26	2	2	
13	26	13	13	
14	26	2	24	
15	26	1	9	
16	26	2	20	
17	26	1	5	
18	26	2	16	
19	26	1	1	$(11 \times 19 = 1 \bmod 26)$
20	26	2	12	
21	26	1	23	
22	26	2	8	
23	26	1	19	
24	26	2	4	
25	26	1	15	
26	26	26	0	

Figure 6. Finding the modular multiplicative inverse.

²¹ The values Alice selects for her encryption key and divisor are not arbitrary. For encryption to work, the divisor must be at least as large as the character set used in the message. From the divisor Eve selects (26), let's assume this character set consists of the lowercase letters of the Latin alphabet, *a* to *z*. As for the public key (11), the only requirement is that its value be *coprime*, or *relatively prime*, with the selected divisor (26). For two numbers to be coprime, the biggest integer that divides both evenly must be 1.

²² Since computers operate on numbers and not letters, pretend the integer 3 represent the letter *c* (because *c* is the 3rd letter of the alphabet).

²³ *Plaintext* and *ciphertext* are the terms of art for unencrypted and encrypted messages, respectively.

The *Keys (e, d)* column contains the set from which Alice selects her public encryption (*e*) and private decryption (*d*) keys. To select her encryption key *e*, Alice need only ensure that the *greatest common divisor (gcd)* of her encryption key *e* and the divisor *n* is 1. This property of *e* ensures that it has a *modular multiplicative inverse* in the set of integers modulo 26.²⁴ Alice selected 11 as her public key *e*, but she could have selected any value between 1 and 26 where $\text{gcd}(e, n) = 1$ (e.g. 1, 3, 5, 7, 9, etc.).

Now that Alice has selected a suitable *e*, she must find its modular multiplicative inverse to serve as her private decryption key *d*. The modular multiplicative inverse of *e* is satisfied by the equation $d \times e = 1 \text{ mod } n$, where *d* is the decryption key, *e* is the encryption key and *n* is the divisor. Looking down the *Inverse? d x e = 1 mod n* column, we find the only value that meets this criterion is 19 ($19 \times 11 = 1 \text{ mod } 26$).²⁵

Now, any message in the range 1 to 26 (or *a* to *z*) that Bob encrypts with Alice's public key *e*, Alice can decrypt with her private key *d*. And having only seen *e* and *n*, and not *d*, Eve cannot break the encryption.

Of course, in this primitive substitution cipher, Eve could very quickly do an exhaustive search of the group of integers modulo 26 to find the inverse of *e*, thereby defeating the encryption.

Textbook RSA

Figure 7 depicts what is referred to in the literature as *textbook* RSA. For our purposes, textbook RSA is close enough to a real-world implementation that we'll conclude our discussion of RSA encryption with it. As with the previous examples, the parameters are set to artificially small values to keep the math manageable.

Parameters						
Step	Alice	Eve	Bob	Public	Private	Calculations
1	5				p	
2	7				q	
3	35	->		n		$5 \times 7 = 35$
4		35	35			
5	24				t	$(5 - 1) \times (7 - 1) = 24$
6	5	->		e		$\text{gcd}(5, 24) = 1$
7		5	5			
8			3			
9		<-	33			$(3^5 = 33 \text{ mod } 35)$
10	33	33				
11	29				d	$(29^5 = 1 \text{ mod } 24)$
12	3					$(33^5 = 3 \text{ mod } 35)$

Figure 7. Encryption and decryption using "textbook" RSA.

Again, Bob wishes to transmit a message to Alice (the letter *c* again, which we encode numerically as 3). In steps 1, 2, 3 and 4, Alice selects two random integers, *p* and *q* (5 and 7), multiplies them to produce a

²⁴ The term *modular multiplicative inverse* sounds scary, but it is to modular exponentiation what division is to multiplication, or logarithm is to exponentiation. Not all members of a set of integers modulo *n*, formally known as a *group*, contain a modular multiplicative inverse; and only an encryption key *e* that has such an inverse will work.

²⁵ Efficient algorithms exist for finding the greatest common divisor between two integers, and for finding the modular multiplicative inverse of an integer in a group of integers. These are, respectively, the Euclidean algorithm and the *extended* Euclidean algorithm. This is important because, if there were not such algorithms, the computational performance of RSA would be too slow for it to have been widely adopted.

divisor n (35), and transmits the divisor to Bob. Eve observes the value of n (35), but not its factors p and q .

Note that in contrast to the previous example, where Alice selected 26 as a divisor (to correspond with the number of letters in the alphabet), she computes the divisor n this time by multiplying two factors, p and q . Alice's only requirement for a suitable n is that its factors p and q both be prime numbers (in the real world these factors would be enormous; more than 600 decimal digits in length!).²⁶

In step 5, Alice applies *Euler's totient function* to the divisor n ; let's call the result of this function t . Given some positive integer n , Euler's totient function tells us the number of positive integers less than n with which n is coprime. For any prime number p , the formula is simple: it is $p - 1$. It should be clear why this is so. If p is prime, we know that the only integers that divide p are 1 and p itself, so every integer in the set 1 to $p - 1$ must be coprime with p . Recall that for two integers to be coprime, the biggest integer that divides both—or their *greatest common divisor*—is 1.

But Alice needs to apply the totient function to the *semiprime* divisor n , which is the product of the two primes p and q (see footnote 26 for the definition of *semiprime*). The formula therefore becomes the product of the totients of n 's factors p and q ; i.e. $(p - 1)(q - 1)$. Plugging in the values from the example we get $(5 - 1)(7 - 1) = 24$.

In plain English, Euler's totient function tells us that there are 24 integers in the set 1 to 35 that are coprime with 35. *Figure 8* depicts this graphically. Indeed, the number of *Yes* entries in the *Coprime?* column confirms the number is 24.

²⁶ The product of the multiplication of any two prime numbers is said to be *semiprime*, because the only numbers that can divide it evenly are 1, the two primes multiplied to produce it (its *factors*), and the product itself; in the present example these numbers are 1, 5, 7 and 35.

e	n	gcd(e, n)	Coprime?
1	35	1	Yes
2	35	1	Yes
3	35	1	Yes
4	35	1	Yes
5	35	5	No
6	35	1	Yes
7	35	7	No
8	35	1	Yes
9	35	1	Yes
10	35	5	No
11	35	1	Yes
12	35	1	Yes
13	35	1	Yes
14	35	7	No
15	35	5	No
16	35	1	Yes
17	35	1	Yes
18	35	1	Yes
19	35	1	Yes
20	35	5	No
21	35	7	No
22	35	1	Yes
23	35	1	Yes
24	35	1	Yes
25	35	5	No
26	35	1	Yes
27	35	1	Yes
28	35	7	No
29	35	1	Yes
30	35	5	No
31	35	1	Yes
32	35	1	Yes
33	35	1	Yes
34	35	1	Yes
35	35	35	No

Figure 8. Euler's totient function.

As with the values of p and q , Alice must keep t secret.

In steps 6 and 7, Alice computes a public key e (5) and transmits it to Bob. Bob will use this key to encrypt the plaintext (3) he transmits to Alice. For e , Alice can select any value in the set 1 to t where the gcd is 1. In the present example, she selects 5 (but 7, 11, 13, 17, 19 or 23 would do).²⁷ Eve observes the value of e .

In steps 8, 9 and 10, Bob encrypts the plaintext (3) and transmits it to Alice. Using the public key e (5) he received from Alice, he raises his plaintext to the power of e , takes the modulus and transmits the resulting ciphertext (33) to Alice ($3^5 \bmod 35 = 33$). Eve observes the ciphertext 33.

In steps 11 and 12, Alice receives and decrypts the ciphertext (33). But in order to decrypt the ciphertext, she must invert it to plaintext. To do this she needs a decryption key d . To find d , Alice computes the modular multiplicative inverse of the public key e (5) in the set of integers modulo t (24).

²⁷ Since the value of e will be used as an exponent in the encryption procedure, it should generally be kept as small as possible to maximize computational performance. All else equal, a small encryption key will not compromise the security of RSA encryption even in a real-world implementation.

The modular multiplicative inverse of 5 in the set of integers modulo 24 is 29; therefore, $d = 29$. Alice raises the ciphertext (33) to the power of d (29), and mods the result with n (35) to retrieve the plaintext ($33^d \bmod 35 = 3$).

To see why the modular multiplicative inverse of 5 in the group of integers modulo 24 is 29, a look at *Figure 9* should help.

Keys (e, d)	Totient t(n)	Divisor gcd(e, t)	Divisor n	Inverse? $d \times e = 1 \bmod t$	Calculations
1	24	1	35	5	(5 x 1 = 5 mod 24)
2	24	2	35	10	(5 x 2 = 10 mod 24)
3	24	3	35	15	(5 x 3 = 15 mod 24)
4	24	4	35	20	(5 x 4 = 20 mod 24)
5	24	1	35	1	(5 x 5 = 1 mod 24)
6	24	6	35	6	(5 x 6 = 6 mod 24)
7	24	1	35	11	...
8	24	8	35	16	
9	24	3	35	21	
10	24	2	35	2	
11	24	1	35	7	
12	24	12	35	12	
13	24	1	35	17	
14	24	2	35	22	
15	24	3	35	3	
16	24	8	35	8	
17	24	1	35	13	
18	24	6	35	18	
19	24	1	35	23	
20	24	4	35	4	
21	24	3	35	9	
22	24	2	35	14	
23	24	1	35	19	
24	24	24	35	0	
25	24	1	35	5	
26	24	2	35	10	
27	24	3	35	15	
28	24	4	35	20	
29	24	1	35	1	(5 x 29 = 1 mod 24)
30	24	6	35	6	
31	24	1	35	11	
32	24	8	35	16	
33	24	3	35	21	
34	24	2	35	2	
35	24	1	35	7	

Figure 9. Finding the modular multiplicative inverse.

The highlighted values in the *Keys (e, d)* column are Alice's public encryption and private decryption keys, e and d , respectively. To select her encryption key e , she finds the first integer greater than one whose gcd with the totient t is 1, or 5 (although she could have selected the equally valid integers 7, 11, 13 and so on, as keys). To select her decryption key d , she finds the modular multiplicative inverse of her encryption key e (5) in the group of integers modulo t (24). The only value that fits the bill is 29, because that is the first d in the equation $d \times e = r \bmod t$ that leaves a remainder r of 1²⁸, and recall that this is the equation that must be satisfied to find an inverse.

The result is that Bob has successfully transmitted an encrypted message to Alice, which Alice has successfully decrypted. But where does this leave Eve? Eve has observed the values of the divisor n , the

²⁸ Actually 5 works, too, but using the same e and d would be a silly choice.

public key e and the ciphertext. But she has not observed the prime factors p and q , and therefore cannot efficiently compute the totient t . Without t , Eve cannot derive the private key d and break the encryption. The reason Eve cannot efficiently compute the totient t is because there is no known, *efficient* way to factor integers. In the present example, this means that given the divisor n (35), there is no efficient way to derive its factors p (5) and q (7).²⁹ This is known as the *integer factorization problem*.

The public-key schemes of secure key exchange and encryption share the property that the computations required for Alice and Bob—finding large primes, computing greatest common divisors and integer exponentiation—all have efficient algorithms; whereas the computations required by Eve to defeat them—the discrete log problem and integer factorization—do not. It is on this basic principle that the efficacy of public-key cryptosystems relies.

RSA Encryption as an Alternative to DH

It could be argued that RSA encryption renders DH key exchange obsolete, because RSA encryption can be used as an alternative to DH for secure key exchange.

To see why this is so, it first helps to realize that, although we refer to DH as a key *exchange* protocol, it is in fact more precisely a key *agreement* protocol. That is, the final product of DH—a symmetric encryption/decryption key to be used in a private message exchange between two parties—is never actually exchanged, but rather computed by both parties independently. Strictly speaking, the encryption key in DH is agreed on, not exchanged.

But with RSA encryption we have the opportunity to *exchange* keys securely.³⁰ Imagine a scenario whereby Alice randomly selects a shared encryption/decryption key, encrypts it with Bob's public encryption key, and transmits the shared key to Bob. Now Alice and Bob both have an identical, shared key, but one that Alice randomly selected and *exchanged* with Bob, rather than one that was computed by Alice and Bob independently. By using RSA to exchange a key, rather than DH to agree on one, Alice and Bob have effectively achieved the same result.

The superiority of one method over the other is open to debate, and beyond the scope of this paper. Nonetheless, the distinction between key agreement and key exchange, and the fact that they can be used interchangeably to serve the same purpose, is useful to be aware of.

Digital Signatures

Digital signatures are the third and final component of the public-key cryptosystem conceived by DH and implemented by RSA. Digital signatures serve three purposes in digital communication: *message integrity*, *authentication of origin* and *non-repudiation of origin*. These properties prove to the recipient

²⁹ Given a small, semiprime divisor n —e.g. 35 in the present example—finding its factors is trivial. For very large values of n , however, factorization is difficult. As always, small values are used in the examples to keep the math simple.

³⁰ A discerning reader might be wondering at this point why one would use public-key encryption to exchange keys for shared-key, or *symmetric*, encryption. After all, if one is already encrypting messages with public keys, and decrypting them with private ones, why bother exchanging symmetric encryption keys? The reason is that public-key encryption is far more expensive computationally than symmetric-key encryption. This means that, in practice, public-key encryption is typically used *only* to exchange symmetric keys; thereafter, all messages between the parties are encrypted using the shared, symmetric key.

of a digitally signed message that it was (a) unaltered in transit, (b) originated from its purported sender and (c) that the purported sender cannot dispute either (a) or (b).³¹

In this scheme, the sender of a message *signs* it using his or her private key, and on receipt of the message the receiver *verifies* it using the signer's public key.³² If the verification fails, this means the private key used to sign the message does not correspond to the public key used to verify it, or that the message was altered after the sender signed it. In short, a digital signature binds the identity of a message's signer to one, and only one, message.

All this comes with a very important caveat: If a private key is stolen from its owner, it can be used to sign messages the owner did not in fact sign. Because of this possibility, it is virtually impossible to prove in a court of law that, just because a message was signed with an owner's private key, the message originated from the owner. In practice, digitally signed messages demand a much weaker burden of proof.³³

Digital signing is the *inverse* of encryption, in the sense that in the latter messages are encrypted with a public encryption key e , and decrypted with a corresponding decryption key d ; whereas in the former messages are *signed* with a decryption key d , and *verified* with a corresponding encryption key e .³⁴

Picking up where *Figure 7. Encryption and decryption using "textbook" RSA* left off, *Figure 10* depicts digital signing. Recall from *Figure 7* that Alice computed a public-private key pair; 5 and 29, respectively. Using the same key pair, Alice now signs a message (4) and transmits it to Bob.³⁵

Step	Alice	Eve	Bob	Calculations
1	4			
2	9	->		$(4^4 = 256 \mod 35)$
3		9	9	
4			4	$(9^5 = 59049 \mod 35)$

Figure 10. Digital signing.

In steps 1, 2 and 3, Alice signs the message (4) and transmits it to Bob. To do this she raises the message (4) to the power of her encryption key d (29), and takes the modulus with the divisor n (35) to compute the signed message (9).

In step 4, Bob verifies the signed message. He raises the signed message (9) to the power of Alice's public key e (5), and arrives back at the original, unsigned message (4).

³¹ Digital signatures are to electronic documents what hand-written signatures are to paper documents; they prove that the signer authorizes the contents of the document. Whereas hand-written signatures can be forged, and/or the documents they belong to altered, neither is possible with digital signatures.

³² Strictly speaking, it is the signature-message *pair* that is verified, not simply the message. A signature is in fact itself just a message; that is, it is a *permutation* of the unsigned message. This permutation is computed by raising the unsigned message to the power of the signer's decryption key modulo the public divisor.

³³ For example, to satisfy oneself that an executable file downloaded from a website can be trusted.

³⁴ Because of this invertibility, encryption and digital signing are said to be *permutations* of one another. Put another way, every message is some other message's ciphertext, and every ciphertext is itself a valid message.

³⁵ For Alice to encrypt the signed message, she would need Bob's public key. Since we already know the mechanism for public-key encryption, these steps are omitted from the diagram so as not to distract from digital signing.

Of course, Eve has observed the signed message (9), and because she also knows Bob's public key (5), she will be able to read it. In a real-world scenario, Alice would have signed the message with her decryption key, and then encrypted it with Bob's public key, before transmitting the message to Bob. On receipt of the message, Bob would perform the inverse procedure; he would decrypt the signed message with his private decryption key, and then verify the signature with Alice's public encryption key.

Man in the Middle Attacks (MITM)

The schemes of the public-key cryptosystem described in this paper are at once powerful and elegant. Moreover, they are virtually ubiquitous in securing electronic commerce and banking on the internet. But a discussion of them would not be complete without pointing out that they are all susceptible to a glaring weakness. This vulnerability is known as a *man-in-the-middle* (MITM) attack.

Imagine the following familiar scenario: Bob wishes to send an encrypted message to Alice. As ever, Eve is listening. Alice computes her public-private key pair and transmits the public key to Bob. Eve, meanwhile, intercepts Alice's key, substitutes it for a public key of her own, and forwards it to Bob. Bob, having no way of knowing that the public key he receives belongs to Eve and not Alice, blithely encrypts the message intended for Alice using Eve's public key. Eve intercepts Bob's message and decrypts it with her private key. Eve has thus broken RSA encryption and, what's more, she has done so without having to resort to sophisticated math.³⁶

This simple example captures the essence of MITM. Any real-world scenario would look very much like some version of it.

Every public-key scheme—whether secure key exchange, encryption, or digital signature—is vulnerable to MITM. How can this be if trillions of US dollars are transacted daily using public-key schemes? The answer is a Frankenstein-like bolt-on called a *public-key infrastructure*, or PKI for short.

Public Key Infrastructure (PKI)

Rivest, Shamir and Adelman, not wanting their legacy to be relegated to one of interesting academic research, founded a company seeking to commercialize their invention not long after publishing *A Method for Obtaining Digital Signatures*. Acutely aware of the MITM problem, the founders knew that *RSA Security Inc.*'s success would depend largely on an effective solution to it. That solution was to be PKI.

At a very high level, a PKI establishes a network of trust between Alice, Bob and a third party known as a *certificate authority* (CA).³⁷ The CA's role in the PKI is to vouch for the authenticity of public keys, which it accomplishes by binding public keys to the identities of their owners. For a PKI to be effective, the CA has to be trusted by both Alice and Bob.

To make this more concrete, take the recent example involving Alice and Bob, where Eve mounted an MITM attack, but this time in the context of a PKI.

³⁶ The MITM attack is an illustration of the *weakest-link* maxim, which in the information security setting holds that a system is only as secure as its weakest link. From the perspective of an attacker, the most rational approach is to attack a system's weakest link. In a public-key cryptosystem, the MITM vulnerability is its weakest link.

³⁷ Alice and Bob being the metaphor for any two parties who wish to participate in a secure message exchange.

After generating her public key (but before transmitting it to Bob), Alice presents it to a CA. The CA vets Alice, and if it concludes she is trustworthy, combines her public key with some name that uniquely identifies Alice into a document known as a *certificate*.³⁸ The CA signs this certificate with its own private key, and returns the signed certificate to Alice. Now, when Bob wants to send an encrypted message to Alice, instead of sending her public key to Bob she sends her CA-signed certificate, which contains her public key. On receipt of the certificate, Bob verifies it using the CA's public key. If the verification succeeds, Bob knows the public key contained in the certificate in fact belongs to Alice, and proceeds to encrypt messages with it. Now, if Eve intercepts Alice's certificate and forwards Bob a phony one of her own, Bob's verification step will fail and he will know he is being attacked.

Though effective at thwarting MITM attacks, PKI is not without weaknesses of its own; chief among them the trustworthiness (or lack thereof) of the CA. Indeed, there are hundreds of CAs the world over, the vast majority of which have sterling reputations. However, CAs themselves are not invulnerable to attack, and in some cases have been compromised.³⁹

Going Forward

For all their power and elegance, cryptosystems based on public-key primitives are imperfect. Nonetheless, they are the best we have in the internet age. Indeed, public-key cryptography is responsible for securing the vast majority of sensitive communication on the internet today, and achieves it with a remarkable degree of success.

The advent of cryptocurrencies, in particular the distributed blockchains on which they run, offers tantalizing prospects for the establishment of trust that today is centralized in a handful of global certificate authorities. A blockchain-based PKI would decentralize trust, spreading it across a vast network of synchronized ledgers, instead of concentrating it in the hands of a vulnerable few. Progress on this front has been halting, not least because its efficacy depends on an alignment of incentives that many would-be applications of a distributed blockchain—such as public-key authentication—lack.⁴⁰

Meanwhile, quantum computing bodes ill for the future of public-key cryptography as we know it. Whereas it would take the most powerful classical computer in the world years to reverse the one-way functions of public-key cryptography, a fit-for-purpose quantum computer could break them in seconds. You can bet that well-funded, state-level actors are working diligently on this problem right now. The invention of such a computer would render classical public-key cryptography instantly obsolete, enabling those who possess it to disrupt the entire system of internet commerce, banking; likely even systems on the which the security of nation-states depends.

³⁸ The official format for public-key certificates is defined by the International Telecommunications Union's (ITU) X.509 standard.

³⁹ In some cases in spectacular fashion. See *DigiNotar* for the quintessential case study.

⁴⁰ This alignment of incentives, known as a *consensus algorithm*, works remarkably well with cryptocurrency, which was the blockchain's original use case. This is because curators (aka *miners*) of such blockchains are compensated with tokens (e.g. bitcoin), which have monetary value, for being honest.